



# Non-linear optimization using MapReduce with SystemML

Deepti Pachauri

(Mentor—Prithviraj Sen, Manager—Rajasekar Krishnamurthy)



Intern Talk

## Overview

- Motivation
- ARIMA
  - ARIMA(1,0,1)
  - ARIMA(p,0,1)
- Optimization
  - Nelder-Mead
  - Distributed Nelder-Mead
  - CG
  - BFGS
- Solver
  - Jacobi
  - GMRES
- Experimental Results

## Motivation

Non-linear optimization is required in many real life problems.

## Motivation

Non-linear optimization is required in many real life problems.

# Time Series Analysis

# Passengers on airport




Arrivals




Departures

# Passengers on airport




Arrivals




Departures

# Passengers on airport



Arrivals



Departures

# Passengers on airport

**Given:** Historical data of passenger movements across the airport.

**Goal:** Predict passenger counts at locations of interest, eg - immigration check points, security lines etc.

**Why:** Proper resource allocation and that makes life easy!!!!!!



# ARIMA: AutoRegressive Integrated Moving Average

Give a time series data  $X_t$

An ARIMA( $p, d, q$ ) process is expressed as

$$(1 - \sum_{i=1}^p \phi_i L^i)(1 - L)^d X_t = (1 + \sum_{i=1}^q \theta_i L^i) \epsilon_t$$

where

$L$  is the **lag operator** given by  $L^i X_t = X_{t-i}$

$\phi_i$  are the autoregressive parameters

$\theta_i$  are the moving average parameters

$\epsilon_t$  are i.i.d. error terms

**Goal: Estimate  $\phi_i$  and  $\theta_i$  for given  $(p, d, q)$**

# Formulation

## ARIMA(1, 0, 1)

$$\hat{X}_2 = \phi_1 X_1 + \theta_1 (X_1 - \hat{X}_1)$$

$$\hat{X}_3 = \phi_1 X_2 + \theta_1 (X_2 - \hat{X}_2)$$

$$\hat{X}_4 = \phi_1 X_3 + \theta_1 (X_3 - \hat{X}_3) \dots \text{and so on}$$

# Formulation

## ARIMA(1, 0, 1)

$$\hat{X}_2 = \phi_1 X_1 + \theta_1 (X_1 - \hat{X}_1)$$

$$\hat{X}_3 = \phi_1 X_2 + \theta_1 (X_2 - \hat{X}_2)$$

$$\hat{X}_4 = \phi_1 X_3 + \theta_1 (X_3 - \hat{X}_3) \dots \text{and so on}$$

$$\begin{pmatrix} 1 & 0 & 0 & \cdots & 0 \\ \theta_1 & 1 & 0 & \cdots & 0 \\ 0 & \theta_1 & 1 & \cdots & 0 \\ \vdots & \ddots & \ddots & & \vdots \\ 0 & 0 & \cdots & \theta_1 & 1 \end{pmatrix} \begin{pmatrix} \hat{X}_1 \\ \hat{X}_2 \\ \hat{X}_3 \\ \vdots \end{pmatrix} = (\theta_1 + \phi_1) \begin{pmatrix} 1 \\ X_1 \\ X_2 \\ \vdots \end{pmatrix}$$

# Formulation

## ARIMA(1, 0, 1)

$$\hat{X}_2 = \phi_1 X_1 + \theta_1 (X_1 - \hat{X}_1)$$

$$\hat{X}_3 = \phi_1 X_2 + \theta_1 (X_2 - \hat{X}_2)$$

$$\hat{X}_4 = \phi_1 X_3 + \theta_1 (X_3 - \hat{X}_3) \dots \text{and so on}$$

For given  $\theta_1$  and  $\phi_1$ , we need to solve

$$A(\theta_1, \phi_1) \hat{X}_t = b(\theta_1, \phi_1)$$

# Formulation

## ARIMA(1, 0, 1)

$$\hat{X}_2 = \phi_1 X_1 + \theta_1 (X_1 - \hat{X}_1)$$

$$\hat{X}_3 = \phi_1 X_2 + \theta_1 (X_2 - \hat{X}_2)$$

$$\hat{X}_4 = \phi_1 X_3 + \theta_1 (X_3 - \hat{X}_3) \dots \text{and so on}$$

For given  $\theta_1$  and  $\phi_1$ , we need to solve

$$A(\theta_1, \phi_1) \hat{X}_t = b(\theta_1, \phi_1)$$

How do we pick correct  $\theta_1$  and  $\phi_1$  for given  $X_t$ ?

$$\min_{\theta_1, \phi_1} \|X_t - \hat{X}_t(\theta_1, \phi_1)\|_2$$

# Formulation

## ARIMA(p, 0, 1)

$$\hat{X}_p = \phi_1 X_{p-1} + \phi_2 X_{p-2} + \phi_3 X_{p-3} + \cdots + \theta_1 (X_{p-1} - \hat{X}_{p-1})$$

$$\begin{pmatrix} 1 & 0 & 0 & \cdots & 0 \\ \theta_1 & 1 & 0 & \cdots & 0 \\ 0 & \theta_1 & 1 & \cdots & 0 \\ \vdots & \ddots & \ddots & & \vdots \\ 0 & 0 & \cdots & \theta_1 & 1 \end{pmatrix} \begin{pmatrix} \hat{X}_1 \\ \hat{X}_2 \\ \hat{X}_3 \\ \vdots \end{pmatrix} = (\theta_1 + \phi_1) \begin{pmatrix} 1 \\ X_1 \\ X_2 \\ \vdots \end{pmatrix} + \phi_2 \begin{pmatrix} 1 \\ 1 \\ X_1 \\ \vdots \end{pmatrix} + \cdots$$

# Formulation

## Summary

$$\begin{aligned} & \min_{\theta, \phi} f(\theta, \phi) \\ \text{subject to } & A(\theta, \phi) \hat{X}_t = b(\theta, \phi) \end{aligned}$$

where  $f(\theta, \phi) = \|X_t - \hat{X}_t(\theta, \phi)\|_2$

- Optimization methods to solve the problem: Nelder-Mead, CG, BFGS, L-BFGS.

# Optimization 1: Nelder-Mead

Downhill simplex method

*Initial point*  $\in \mathbb{R}^n$






# Optimization 1: Nelder-Mead

## Downhill simplex method


*Simplex*  $n + 1$  points



# Optimization 1: Nelder-Mead

## Downhill simplex method


*Simplex*  $n + 1$  points



# Optimization 1: Nelder-Mead

Downhill simplex method


*Reflection*



# Optimization 1: Nelder-Mead

Downhill simplex method


*Expand*



# Optimization 1: Nelder-Mead

Downhill simplex method


*Contraction*




# Optimization 1: Nelder-Mead

Downhill simplex method


*Shrink*




# Optimization 1: **Distributed** Nelder-Mead




**SystemML:** Distribute simplex on multiple machines.



Machine 1




Machine 2




Machine 3

...


# Optimization 1: **Distributed** Nelder-Mead




**SystemML:** Distribute simplex on multiple machines.



Machine 1



Machine 2




Machine 3

...



## Optimization 1: **Distributed** Nelder-Mead



**SystemML:** Distribute simplex on multiple machines.

- Collect “local best” at each machine.
- Identify “global best”
- Communicate “global best” back to each machine
- Proceed with Nelder-Mead steps  
i.e., *reflection, expansion, contraction*
- If no update recorded anywhere, *shrink* simplex on each machine

## Optimization 2: Conjugate Gradient

### First-order optimization method

Generate a sequence of points  $(\theta^0, \phi^0), (\theta^1, \phi^1), (\theta^2, \phi^2), \dots$  such that

$$f(\theta^0, \phi^0) \geq f(\theta^1, \phi^1) \geq f(\theta^2, \phi^2), \dots$$

One possible way to obtain such a sequence is – move along a descent direction of the function  $\Delta(\theta, \phi)$

$$(\theta^{k+1}, \phi^{k+1}) := (\theta^k, \phi^k) + t_k \Delta(\theta^k, \phi^k)$$

# Optimization 2: Conjugate Gradient

## Algorithm

**given** initial point  $(\theta, \phi)$ .

**repeat**

- Determine descent direction  $\Delta(\theta, \phi)$
- Choose a step size  $t > 0$ .
- *Update*  $(\theta, \phi) := (\theta, \phi) + t\Delta(\theta, \phi)$ .

**until** stopping criterion is satisfied.

## Optimization 2: Conjugate Gradient

### Algorithm

**given** initial point  $(\theta, \phi)$ .

**repeat**

- **Set**  $\Delta(\theta, \phi) = -\nabla f$ .  
Use finite difference method to compute  $\nabla f$ .
- Choose a step size  $t > 0$ .  
Backtracking line search  
 $t := 1, \alpha \in (0, 0.5), \beta \in (0, 1)$   
**while**  
 $f((\theta, \phi) + t\Delta(\theta, \phi)) > f(\theta, \phi) + \alpha t \nabla f^\top \Delta(\theta, \phi), \quad t := \beta t$
- **Update**  $(\theta, \phi) := (\theta, \phi) + t\Delta(\theta, \phi)$ .

**until**  $\|\nabla f\|_2 \leq \eta$

# Optimization 3

## Newton method

Use second order Taylor expansion to obtain the descent direction

$$\Delta(\theta, \phi) = -[Hf]^{-1} \nabla f$$

$Hf$  is the Hessian matrix.

# Optimization 3

## Newton method

Use second order Taylor expansion to obtain the descent direction

$$\Delta(\theta, \phi) = -[Hf]^{-1} \nabla f$$

$Hf$  is the Hessian matrix. **Prohibitively Expensive**

## Optimization 3: **Broyden–Fletcher–Goldfarb–Shanno**

### Quasi-Newton method

- Approximate Hessian matrix at each iteration by constructing a rank-two update matrix using  $\nabla f$ .
- Efficiently use Sherman–Morrison formula to obtain inverse of the approximate Hessian.

Note: approximate Hessian matrix is denoted by  $B$ .

# Optimization 3: Broyden–Fletcher–Goldfarb–Shanno

## Algorithm

**given** initial point  $(\theta_0, \phi_0)$  and  $B_0 = I$ .

**repeat**

- Determine descent direction  $\Delta(\theta_k, \phi_k) = -B_k^{-1} \nabla f$
- Choose a step size  $t_k > 0$   
Update  $(\theta_{k+1}, \phi_{k+1}) = (\theta_k, \phi_k) + t_k \Delta(\theta_k, \phi_k)$
- Set  $s_k = t_k \Delta(\theta_k, \phi_k)$
- Calculate  $y_k = \nabla f(\theta_{k+1}, \phi_{k+1}) - \nabla f(\theta_k, \phi_k)$
- $B_{k+1}^{-1} = B_k^{-1} + \frac{(s_k^\top y_k + y_k^\top B_k^{-1} y_k)(s_k s_k^\top)}{(s_k^\top y_k)^2} - \frac{B_k^{-1} y_k s_k^\top + s_k y_k^\top B_k^{-1}}{s_k^\top y_k}$

**until** stopping criterion is satisfied



# Optimization: How do we choose??

## Nelder-Mead

**Pros:** Simple, no derivative required, good for non-convex problems.

**Cons:** Local search method – can easily get stuck in local minimum, too many function evaluations required, very slow.

## Conjugate Gradient

**Pros:** Simple.

**Cons:** Too slow near minimum, ill-defined for non-differentiable functions.

## BFGS

**Pros:** **Works in most cases**, fast.

**Cons:** ill-defined for non-differentiable functions, computationally expensive (**try L-BFGS**), do not necessarily converge.

# Revisit – Formulation

## Summary

$$\begin{aligned} & \min_{\theta, \phi} f(\theta, \phi) \\ & \text{subject to } A(\theta, \phi) \hat{X}_t = b \end{aligned}$$

where  $f(\theta, \phi) = \|X_t - \hat{X}_t(\theta, \phi)\|_2$

**Note:**  $A \in \mathbb{R}^{T \times T}$  and  $T$  is large for large time-series.

Solve non-symmetric sparse system of linear equations efficiently?

# Solver 1: Jacobi method

**Key idea:**  $A = D + R$

$$\begin{aligned}(D + R)\hat{X} &= b \\ D\hat{X} &= b - R\hat{X}\end{aligned}$$

Here  $D$  is a diagonal matrix with  $D_{ii} = A_{ii}$ .  $R$  constitute off-diagonal entries of  $A$ .

**Iterate until convergence:**  $\hat{X}^{k+1} = D^{-1}b - D^{-1}R\hat{X}^k$

# Solver 1: Jacobi method

**Key idea:**  $A = D + R$

$$\begin{aligned}(D + R)\hat{X} &= b \\ D\hat{X} &= b - R\hat{X}\end{aligned}$$

Here  $D$  is a diagonal matrix with  $D_{ii} = A_{ii}$ .  $R$  constitute off-diagonal entries of  $A$ .

**Iterate until convergence:**  $\hat{X}^{k+1} = D^{-1}b - D^{-1}R\hat{X}^k$

For problem at hand,

$$\begin{pmatrix} 1 & 0 & 0 & \cdots & 0 \\ \theta_1 & 1 & 0 & \cdots & 0 \\ 0 & \theta_1 & 1 & \cdots & 0 \\ \vdots & \ddots & \ddots & & \vdots \\ 0 & 0 & \cdots & \theta_1 & 1 \end{pmatrix} \begin{pmatrix} \hat{X}_1 \\ \hat{X}_2 \\ \hat{X}_3 \\ \vdots \end{pmatrix} = (\theta_1 + \phi_1) \begin{pmatrix} 1 \\ X_1 \\ X_2 \\ \vdots \end{pmatrix}$$

# Solver 1: Jacobi method

**Key idea:**  $A = D + R$

$$\begin{aligned}(D + R)\hat{X} &= b \\ D\hat{X} &= b - R\hat{X}\end{aligned}$$

Here  $D$  is a diagonal matrix with  $D_{ii} = A_{ii}$ .  $R$  constitute off-diagonal entries of  $A$ .

**Iterate until convergence:**  $\hat{X}^{k+1} = D^{-1}b - D^{-1}R\hat{X}^k$

For problem at hand,

$$D = \begin{pmatrix} 1 & 0 & 0 & \cdots & 0 \\ 0 & 1 & 0 & \cdots & 0 \\ 0 & 0 & 1 & \cdots & 0 \\ \vdots & \ddots & \ddots & & \vdots \\ 0 & 0 & \cdots & 0 & 1 \end{pmatrix}; \quad R = \begin{pmatrix} 0 & 0 & 0 & \cdots & 0 \\ \theta_1 & 0 & 0 & \cdots & 0 \\ 0 & \theta_1 & 0 & \cdots & 0 \\ \vdots & \ddots & \ddots & & \vdots \\ 0 & 0 & \cdots & \theta_1 & 0 \end{pmatrix}$$

# Solver 1: Jacobi method

**Key idea:**  $A = D + R$

$$\begin{aligned}(D + R)\hat{X} &= b \\ D\hat{X} &= b - R\hat{X}\end{aligned}$$

Here  $D$  is a diagonal matrix with  $D_{ii} = A_{ii}$ .  $R$  constitute off-diagonal entries of  $A$ .

**Iterate until convergence:**  $\hat{X}^{k+1} = D^{-1}b - D^{-1}R\hat{X}^k$

This method is guaranteed to converge given  
**diagonal dominance** of  $A$

## Solver 2: Generalized minimal residual method

**Key idea:** *Krylov subspace*

$$\begin{aligned} & k\text{-th Krylov subspace of } A\hat{X} = b \\ & \mathcal{K}_k = \text{span}(b, Ab, A^2b, \dots, A^{k-1}b) \end{aligned}$$

- GMRES approximate the exact solution of  $A\hat{X} = b$  by  $\hat{X}^k \in \mathcal{K}_k$
- $\hat{X}^k \in \mathcal{K}_k$  is the vector that minimizes residual

$$r_k = b - A\hat{X}^k$$

## Solver 2: Generalized minimal residual method

---

### Algorithm 1 Arnoldi iteration for $Q_k$

---

**Require:**  $A, b$

**Compute**  $q_1 = \frac{b}{||b||_2}$

**for**  $i = 1$  to  $k$  **do**

$v = Aq_i$

**for**  $j = 1$  to  $i$  **do**

$S(j, i) = v^\top q_j$

$v = v - S(j, i) * q_j$

**end for**

$q_{i+1} = \frac{v}{||v||_2}$

$S(i + 1, i) = ||v||_2$

**end for**

**Ensure:** Orthonormal basis  $Q_k$  and upper *Hessenberg* matrix  $S_k$

---



## Solver 2: Generalized minimal residual method

Solution for  $A\hat{X} = b$

$$e_1 = (1, 0, 0, \dots, 0)$$

$$e_1 = \|b\|_2 e_1$$

$$y_k = S_k(1:k, 1:k) \setminus e_1 \quad \text{We used exact solver here}$$

$$\hat{X}^k = Q_k y_k \quad \text{Best solution in } \mathcal{K}_k$$

Our experimental results used **GMRES**

# Preliminary results

We present preliminary results with

- Optimization routine – BFGS
- Solver – GMRES ( $K_{10}$ )
- Line search – backtracking ( $\alpha = 0.0001, \beta = 0.9$ )

# Preliminary results: ARIMA (1, 0, 1)

**Time step–1 hour**

X1011

	SystemML	R
$\phi_1$	-0.7233	NaN
$\theta_1$	0.6932	NaN

X1157

	SystemML	R
$\phi_1$	0.8662	$0.8662 \pm 0.0189$
$\theta_1$	-0.3316	$-0.3318 \pm 0.0417$

# Preliminary results: ARIMA (1, 0, 1)

**Time step–1 hour**

X1158

	SystemML	R
$\phi_1$	0.9053	$0.9976 \pm 0.0017$
$\theta_1$	-0.5076	$-0.014 \pm 0.0246$

X1178

	SystemML	R
$\phi_1$	0.9479	$1.0001 \pm 0.0006$
$\theta_1$	-0.8860	$-0.9944 \pm 0.0026$

# Preliminary results: ARIMA(5, 0, 1)

**Time step—1 hour**

X1157

	SystemML	R
$\phi_1$	0.8920	$0.892 \pm 0.0759$
$\phi_2$	-0.3264	$-0.3265 \pm 0.0562$
$\phi_3$	0.4233	$0.4233 \pm 0.0313$
$\phi_4$	-0.3277	$-0.3278 \pm 0.0399$
$\phi_5$	0.2202	$0.2203 \pm 0.0242$
$\theta_1$	-0.2133	$-0.2133 \pm 0.0756$

# Preliminary results: More

## X1011 Time step–15mins

ARIMA(1, 0, 1)

	SystemML	R
$\phi_1$	0.2318	$0.2315 \pm 0.0701$
$\theta_1$	-0.0285	$-0.028 \pm 0.0726$

ARIMA(3, 0, 1)

	SystemML	R
$\phi_1$	-0.1549	$-0.1381 \pm 0.2211$
$\phi_2$	0.0691	$0.0657 \pm 0.0472$
$\phi_3$	0.0391	$0.0394 \pm 0.0122$
$\theta_1$	0.3591	$0.3423 \pm 0.2212$

# Summary

- Implemented Serial Nelder-Mead in DML.
- Implemented Distributed Nelder-Mead in DML.
- Implemented CG in DML.
- Implemented BFGS in DML.
- Implemented L-BFGS in DML.
- Implemented GMRES in DML.
- Made useful observations to improve SystemML.

Thank You!





# Preliminary results: X1157

**Time step—1 hour**

ARIMA(2, 0, 1)

	SystemML	R
$\phi_1$	0.0442	$0.0184 \pm 0.0459$
$\phi_2$	0.4410	$0.4636 \pm 0.0389$
$\theta_1$	0.7626	$0.787 \pm 0.0367$

ARIMA(3, 0, 1)

	SystemML	R
$\phi_1$	0.2588	$0.2585 \pm 0.0479$
$\phi_2$	0.1320	$0.1323 \pm 0.0398$
$\phi_3$	0.3413	$0.3413 \pm 0.0228$
$\theta_1$	0.4140	$0.4143 \pm 0.0481$